

LLM-Augmented Log Anomaly Detection: Automated Root Cause Narration for AIOps using Lightweight Foundation Models on Real-World Campus Infrastructure

Trinh Hoang Tu¹ and Cao Tien Thanh^{1*}

Department of Cybersecurity, Faculty of Information Technology,
Ho Chi Minh City University of Foreign Languages - Information Technology
(HUFLIT),
Ho Chi Minh City, Vietnam
tht.csec2005@gmail.com, thanhct@huflit.edu.vn

Abstract. Modern enterprise campus networks generate massive volumes of heterogeneous, unstructured log streams across multi-site infrastructures. While deep learning models achieve high accuracy in sequence anomaly detection, traditional AIOps platforms suffer from a critical explainability bottleneck: raw reconstruction error scores fail to communicate actionable root causes, forcing Security Operations Center (SOC) analysts to manually triage thousands of log messages under severe time constraints. In this paper, we propose a novel **LLM-Augmented AIOps Framework** synthesizing streaming Drain3 log template parsing, a hybrid Temporal Convolutional Network and Transformer (TCN-Transformer) autoencoder, and a lightweight quantized Foundation Model (Phi-3 / Mistral via INT4 LoRA fine-tuning) for real-time automated Root Cause Analysis (RCA) narration. Evaluated on a production 19.28 GB enterprise campus log dataset from HUFLIT alongside standard benchmarks (BGL, HDFS, UNSW-NB15), our framework achieves a detection F1-score of 0.9521, ROC-AUC of 0.9784, reduces average incident triage time by > 99% (125.18 ms LLM narration latency vs 15.4 minutes human manual inspection), and delivers high-fidelity ROUGE-L score (0.8420) compared to expert ground-truth playbooks. All source code, dataset scripts, and benchmark evaluation artifacts are publicly released.

Keywords: AIOps · Log Anomaly Detection · Foundation Models · Root Cause Analysis · TCN-Transformer · Drain3 · Enterprise Campus Infrastructure · SOICT 2026.

1 Introduction

Higher education and enterprise ICT infrastructures produce high-density syslog telemetry across firewalls, active directory domain controllers, learning management systems (LMS), and core routing backbones. Rapid anomaly identification

* Corresponding author: thanhct@huflit.edu.vn

is essential for preventing lateral threat movement, credential theft, and network outages.

Recent advancements in deep sequence modeling—such as DeepLog [2], LogAnomaly [3], and Transformer Autoencoders—have dramatically improved anomaly detection accuracy. However, existing approaches present a significant operational bottleneck:

1. **Explainability Deficit:** Anomaly detectors output numeric reconstruction errors or probabilities (e.g., $RE = 0.942$), which fail to explain *why* an anomaly occurred or *what* specific attack vector is active.
2. **Prolonged Human Triage:** SOC analysts must manually search raw log repositories, correlate multi-service timestamps, and draft mitigation steps, leading to Mean Time to Respond (MTTR) exceeding 15 minutes per incident.

To bridge this gap, we present an **LLM-Augmented AIOps Framework** engineered specifically for enterprise campus environments. By coupling Drain3 streaming template extraction and TCN-Transformer sequence modeling with a lightweight INT4-quantized Large Language Model (LLM), our architecture automatically generates natural language Root Cause Analysis (RCA) narratives and structured mitigation playbooks within milliseconds of anomaly detection.

The key contributions of this paper include:

- **End-to-End AIOps Architecture:** A multi-stage pipeline combining Drain3 template extraction, TCN-Transformer sequence anomaly detection, and RAG-guided LLM narration.
- **Lightweight Foundation Model RCA:** Fine-tuning quantized Phi-3 / Mistral models to produce structured JSON incident summaries, attack vector explanations, and remediation steps.
- **Evaluation on 19.28 GB Real-World Production Campus Logs:** Rigorous empirical benchmarking using a 19.28 GB enterprise campus network log dataset collected from HUFLIT alongside BGL, HDFS, and UNSW-NB15 benchmarks.

2 Related Work

Log anomaly detection has evolved from statistical rule-based techniques to deep sequential models. Table 1 summarizes the state-of-the-art architectures in log intelligence.

While existing systems focus strictly on computing anomaly scores, our framework incorporates an LLM narration layer to convert raw sequence vectors into human-readable SOC incident playbooks.

3 System Architecture & Methodology

The overall workflow of the proposed system is depicted in Fig. 1.

Table 1. Taxonomy and Comparison of Existing Log Anomaly Detection Systems

System / Approach	Parsing Method	Core Detector	Automated RCA	Inference Latency
PCA / Invariant Mining	Fixed Rules	Linear Projection	No	< 5 ms
DeepLog [2]	Spell	Stacked LSTM	No	≈ 24 ms
LogAnomaly [3]	FT-Tree	FT-LSTM + Word2Vec	Partial	≈ 31 ms
SwissLog	Drain3	Attentive Bi-LSTM	No	≈ 18 ms
Ours (LLM-Augmented)	Drain3 Stream	TCN-Transformer	Yes (LLM Narrator)	0.0128 ms + 125 ms LLM

**Fig. 1.** End-to-End LLM-Augmented AIOps Pipeline: Streaming Drain3 Parsing, TCN-Transformer Anomaly Detection, and Foundation Model RCA Narration.

3.1 Streaming Drain3 Log Parsing

Raw log messages $L = \{l_1, l_2, \dots, l_n\}$ are ingested in real-time. Drain3 [1] uses a fixed-depth parse tree to sanitize variable tokens (IP addresses, timestamps, numerical IDs) into structured template identifiers $T_i \in \mathcal{V}_{\text{vocab}}$. Token similarity between sequence a and b is computed as:

$$\text{Sim}(a, b) = \frac{\sum_{j=1}^{|a|} \mathbb{I}(a_j = b_j)}{\max(|a|, 1)} \quad (1)$$

3.2 TCN-Transformer Sequence Autoencoder

Parsed template sequences $S = [T_1, T_2, \dots, T_k]$ are embedded into vector representations $E \in \mathbb{R}^{k \times d}$. 1D dilated causal convolutions extract short-term temporal dependencies:

$$y(t) = (x *_d f)(t) = \sum_{i=0}^{k-1} f(i) \cdot x(t - d \cdot i) \quad (2)$$

A multi-head Transformer Encoder processes the TCN output H_{TCN} using Multi-Head Self-Attention:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (3)$$

Reconstruction loss over template vocabulary $\mathcal{V}$ yields the sequence anomaly score S_{anomaly} :

$$S_{\text{anomaly}} = \frac{1}{k} \sum_{i=1}^k \mathcal{L}_{\text{CE}}(\hat{y}_i, y_i) \quad (4)$$

3.3 Foundation Model RCA Narration Engine

When $S_{\text{anomaly}} > \tau$, the anomalous log window, template context, and metadata are formatted into a structured prompt $\mathcal{P}$. A quantized Phi-3 / Mistral LLM generates a JSON response containing:

1. **Severity Level:** {CRITICAL, HIGH, MEDIUM, LOW}
2. **Root Cause Summary:** Concise description of the underlying fault or attack.
3. **Remediation Playbook:** Ordered list of immediate counter-measures.

4 Experimental Evaluation

4.1 Dataset Specification

We evaluated our framework on four distinct datasets detailed in Table 2.

Table 2. Benchmark Dataset Statistics

Dataset	Raw File Size	Total Log Lines	Environment
HUFLIT Campus Logs	19.28 GB	14,850,220	Real-World Enterprise
BGL (BlueGene/L)	744 MB	4,747,963	Supercomputer
HDFS	1.58 GB	11,175,629	Distributed Cluster
UNSW-NB15	100 MB	254,004	Network Intrusion

4.2 Detection Accuracy Comparison

Table 3 benchmarks TCN-Transformer against baseline log detection models. Our hybrid architecture achieves optimal F1-score (0.9521) with low inference latency (0.0128 ms per sequence).

Table 3. Anomaly Detection Accuracy on HUFLIT Campus Logs

Model	Precision	Recall	F1-Score	ROC-AUC
PCA	0.7420	0.6850	0.7123	0.7310
DeepLog (LSTM) [2]	0.8850	0.9120	0.8983	0.9150
LogAnomaly [3]	0.9010	0.9240	0.9123	0.9320
TCN-Transformer (Ours)	0.9482	0.9561	0.9521	0.9784




Fig. 2. Reconstruction Loss Convergence.




Fig. 3. ROC and PR Curves across Datasets.

4.3 Training Loss and ROC/PR Curve Analysis

Fig. 2 and Fig. 3 demonstrate training loss convergence and ROC/Precision-Recall curves across evaluation datasets.

4.4 LLM RCA Narration Quality & Triage Efficiency

Evaluating LLM-generated RCA narratives against expert-annotated incident playbooks yields an average ROUGE-L score of 0.8420 and BLEU-4 of 0.7815. Fig. 4 shows that automated narration reduces human incident triage time from 15.4 minutes to 125.18 ms per incident ($> 99\%$ latency reduction).




Fig. 4. Incident Triage Latency Comparison (Log Scale): Manual Inspection vs LLM Narration.

5 Real-World Campus Case Studies

5.1 Case Study 1: Active Directory Credential Stuffing

A sequence of 45 failed authentication attempts within 3.2 seconds on Active Directory triggers an anomaly score $S_{\text{anomaly}} = 0.942$. The LLM engine generates:

Severity: CRITICAL

Root Cause: High-density Credential Stuffing Attack targeting LDAP directory.

Playbook: 1. Push OpenFlow drop rule to Core Switch for source IP; 2. Enforce temporary 15-min account lockout; 3. Trigger 2FA prompt.

5.2 Case Study 2: Inter-VLAN STP Loop Broadcast Storm

An abnormal surge of broadcast packet log templates on Edge Switch 4 causes high packet drop. The LLM engine pinpoints:

Severity: HIGH

Root Cause: Spanning Tree Protocol (STP) loop causing link saturation.

Playbook: 1. Enable BPDU Guard on Switch Port 12; 2. Shift academic VLAN traffic to backup optical link.

6 Threats to Validity & Operational Feasibility

Internal validity is mitigated by evaluating on real production syslog data alongside public benchmarks. External validity is supported by testing on multi-domain environments (campus network, supercomputer, distributed cluster, IoT). Quantizing the LLM to INT4 ensures memory footprints under 2.4 GB, permitting deployment on edge gateway servers.

7 Conclusion

This paper presented an LLM-Augmented AIOps framework for log anomaly detection and automated Root Cause Analysis on real-world enterprise campus infrastructure. By integrating Drain3 streaming parsing, TCN-Transformer sequence modeling, and lightweight quantized LLM narration, the system achieves sub-second explainable incident triage ($F1 = 0.9521$, RCA latency = 125.18 ms) evaluated on 19.28 GB production campus logs.

Acknowledgment

The authors thank the Faculty of Information Technology at HUFLIT for providing infrastructure logs and academic support. SOICT 2026 is acknowledged for peer-review management.

References

1. He, P., Zhuang, J., Ma, X., et al.: Drain: An Online Log Parsing Approach Based on Fixed Depth Tree. In: ICWS (2017).
2. Du, M., Li, F., Zheng, G., Srikumar, V.: DeepLog: Anomaly Detection and Diagnosis from System Logs Through Deep Learning. In: ACM CCS (2017).
3. Meng, W., Liu, Y., Zhu, Y., et al.: LogAnomaly: Unsupervised Detection of Sequential and Quantitative Anomalies in Unstructured Logs. In: IJCAI (2019).
4. Bai, S., Kolter, J.Z., Koltun, V.: An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling. arXiv:1803.01271 (2018).
5. Abbadin, M., et al.: Phi-3 Technical Report: Highly Capable Language Model Locally on Phone. arXiv:2404.14219 (2024).