

LLM-Augmented Log Anomaly Detection: Automated Root Cause Narration for AIOps using Lightweight Foundation Models on Real-World Campus Infrastructure

Trinh Hoang Tu¹ and Cao Tien Thanh^{1*}

Faculty of Information Technology,
Ho Chi Minh City University of Foreign Languages - Information Technology,
Ho Chi Minh City, Vietnam
tth.csec2005@gmail.com, thanhct@huflit.edu.vn

Abstract. Modern enterprise and educational campus infrastructures generate massive volumes of heterogeneous, unstructured log streams across multi-site networks. While deep learning models achieve high statistical accuracy in sequence anomaly detection, traditional Artificial Intelligence for IT Operations (AIOps) platforms suffer from a critical explainability bottleneck: raw reconstruction error scores fail to communicate actionable root causes, forcing Security Operations Center (SOC) analysts to manually triage thousands of log messages under severe time constraints. In this paper, we propose a novel **LLM-Augmented AIOps Framework** synthesizing streaming Drain3 log template parsing, a hybrid Temporal Convolutional Network and Transformer (TCN-Transformer) autoencoder, and a lightweight quantized Foundation Model (Phi-3 / Mistral via INT4 quantization and RAG-guided prompting) for real-time automated Root Cause Analysis (RCA) narration. Evaluated on a production 19.28 GB enterprise campus log dataset (14,850,220 log lines) from HUFLIT alongside standard benchmarks (BGL, HDFS, and UNSW-NB15), our hybrid detector achieves an anomaly detection F1-score of 0.9521 and ROC-AUC of 0.9784 on campus logs (outperforming DeepLog and LogAnomaly baselines), while maintaining a sub-millisecond sequence scoring latency of 0.0128 ms. Furthermore, the quantized foundation model reduces average incident triage time by > 99% (125.18 ms LLM narration latency vs 15.4 minutes human manual triage) and delivers high semantic fidelity (ROUGE-L of 0.8420, BLEU-4 of 0.7815) against expert-curated ground-truth playbooks. All source code, dataset parsing pipelines, and benchmark artifacts are prepared for open-source release upon publication.

Keywords: AIOps · Log Anomaly Detection · Large Language Models · Root Cause Analysis · TCN-Transformer Autoencoder · Drain3 · Enterprise Campus Infrastructure.

* Corresponding author: thanhct@huflit.edu.vn

1 Introduction

Enterprise networks and higher education IT infrastructures generate high-density syslog telemetry continuously streaming from perimeter firewalls, Active Directory domain controllers, Learning Management Systems (LMS), student authentication portals, and core routing backbones [1, 2]. Rapid and accurate anomaly detection within these log streams is paramount for thwarting lateral cyber attacks, credential brute-forcing, data exfiltration, and service disruptions [3].

Recent advancements in deep learning have introduced sequence-based log anomaly detection models, notably DeepLog (LSTM-based next-token prediction) [2], LogAnomaly (semantic template vectors) [3], and Transformer-based autoencoders [5, 4]. While these architectures achieve strong empirical detection fidelity on synthetic or curated benchmarks, existing AIOps paradigms exhibit two fundamental operational bottlenecks in production environments:

1. **Explainability Deficit:** Contemporary anomaly detectors produce scalar anomaly scores or probabilistic reconstruction errors (e.g., $RE = 0.942$). These numerical scores provide zero contextual explanation regarding *why* a sequence is anomalous, *what* specific attack vector or hardware malfunction is occurring, or *which* downstream services are impacted.
2. **Prolonged Incident Triage & SOC Alert Fatigue:** Security Operations Center (SOC) analysts and network reliability engineers must manually query raw distributed log stores, correlate timestamps across multi-tier appliances, and formulate remediation scripts. In enterprise campus settings, this manual triage process frequently requires 15–30 minutes per critical incident, allowing fast-propagating threats (such as automated credential stuffing or spanning tree broadcast loops) to cause widespread disruption.

To overcome these limitations, we propose an end-to-end **LLM-Augmented AIOps Framework** tailored for real-world enterprise campus networks. As illustrated in Fig. 1, our framework couples streaming log parsing and deep sequence anomaly detection with a lightweight, quantized Foundation Model to automatically produce natural language Root Cause Analysis (RCA) narratives and structured incident mitigation playbooks within milliseconds of anomaly detection.

The core research contributions of this paper are summarized as follows:

- **End-to-End AIOps Architecture:** We construct a unified 3-stage pipeline integrating online Drain3 template extraction, a hybrid TCN-Transformer autoencoder for microsecond-level sequence anomaly scoring, and an LLM-driven narration engine.
- **Hybrid TCN-Transformer Autoencoder:** We design a sequential deep autoencoder combining 1D dilated causal convolutions (capturing local short-term temporal dependencies) with Multi-Head Self-Attention (capturing long-range contextual correlations across log templates), achieving superior detection accuracy with sub-millisecond inference (0.0128 ms per sequence).




Fig. 1. Architectural Overview of the LLM-Augmented AIOps Framework: (1) Streaming Drain3 log parsing and template encoding; (2) TCN-Transformer autoencoder sequence scoring; and (3) Quantized Foundation Model Root Cause Narration Engine.

- **Lightweight Foundation Model RCA Narration:** We formulate a Retrieval-Augmented Generation (RAG) prompting architecture using INT4-quantized open-source foundation models (Microsoft Phi-3 Mini / Mistral-7B) to synthesize structured, human-readable RCA narratives and step-by-step mitigation playbooks.
- **Empirical Validation on 19.28 GB Real-World Production Logs:** We validate our framework on a comprehensive 19.28 GB enterprise campus log dataset (14,850,220 log lines) collected from the HUFLIT production infrastructure, alongside standard public benchmarks (BGL, HDFS, and UNSW-NB15), demonstrating an anomaly detection F1-score of 0.9521, ROC-AUC of 0.9784, and a $> 99\%$ reduction in SOC triage latency.

2 Related Work

2.1 Log Parsing and Template Extraction

Raw log messages are unstructured strings containing static message templates interleaved with dynamic runtime variables (e.g., IP addresses, user IDs, timestamps). Log parsing converts raw log lines into discrete template identifiers. Early approaches utilized regular expressions or invariant mining [1]. He et al. [1] proposed Drain, an online parser using a fixed-depth parse tree that achieves high parsing accuracy and throughput. Drain3 extends Drain with persistent caching and streaming support. In this work, we employ Drain3 to sanitize high-throughput syslog streams in real time.

2.2 Deep Learning for Log Anomaly Detection

DeepLog [2] pioneered the application of deep sequence modeling to log analytics by training a Long Short-Term Memory (LSTM) network to predict the next log template in nominal execution paths. LogAnomaly [3] introduced template2vec embeddings to model sequential and quantitative log patterns simultaneously. LogRobust [4] applied Bi-LSTMs with attention to handle unstable and unseen log templates. Recently, Transformer architectures [7, 5] have been adopted to

log intelligence, leveraging self-attention mechanisms. However, existing detectors focus exclusively on binary classification and numerical error computation, lacking semantic root cause explanation.

2.3 Foundation Models & LLMs in AIOps

Large Language Models (LLMs) have demonstrated exceptional reasoning and summarization capabilities across software engineering and system telemetry [8, 10, 11]. Recent exploratory studies have investigated LLMs for cloud incident management and root cause analysis [12, 13]. Nonetheless, deploying massive billion-parameter models in real-time operational environments is cost-prohibitive. Our work addresses this by combining INT4 post-training quantization with structured RAG prompting on specialized campus log sequences. Table 1 compares our framework against representative state-of-the-art systems.

Table 1. Taxonomy and Comparison with Existing Log Anomaly Detection Frameworks

Framework	Parsing Method	Sequence Detector	Automated RCA	Explainability Output	Inference Latency
PFA / Invariants [1]	Fixed Regex	Linear Projection	No	None	< 5.0 ms
DeepLog [2]	Spell / Drain	Stacked LSTM	No	Anomaly Score	≈ 24.5 ms
LogAnomaly [3]	FT-Tree	FT-LSTM + Template2Vec	Partial	Anomaly Score	≈ 31.2 ms
SwissLog [5]	Drain3	Attentive Bi-LSTM	No	Attention Weights	≈ 18.4 ms
Ours (LLM-Augmented)	Drain3 Stream	TCN-Transformer	Yes (LLM Narrator)	JSON Incident Playbook	0.0128 ms + 125 ms LLM

3 System Architecture & Methodology

3.1 Architectural Workflow

The proposed LLM-Augmented AIOps Framework operates across three pipelined stages: (i) **Stage 1 (Online Parsing)**: Raw syslog streams are ingested and tokenized by Drain3 into structured template IDs; (ii) **Stage 2 (Sequence Anomaly Detection)**: Sliding template windows are vectorized and scored by a hybrid TCN-Transformer Autoencoder; (iii) **Stage 3 (Foundation Model RCA Narration)**: If a sequence’s reconstruction error exceeds a calibrated anomaly threshold τ , an enriched incident prompt is compiled and dispatched to an INT4-quantized LLM to generate an actionable incident RCA summary and remediation playbook.

3.2 Streaming Drain3 Log Template Parsing

Given a continuous log stream $\mathcal{L} = \{l_1, l_2, \dots, l_N\}$, each message l_i is preprocessed to mask dynamic tokens (e.g., IPv4/IPv6 addresses, MAC addresses, hex memory addresses, timestamps, numerical transaction IDs) using regular expression sanitizers. Drain3 [1] processes the sanitized message through a fixed-depth parse tree. The root node branches on log message length (number of tokens), subsequent internal nodes branch on token positions, and leaf nodes maintain clusters of log templates:

$$\mathcal{T}_j = \langle \text{Template String, Template ID } k \rangle \quad (1)$$

The token similarity between an incoming message token list s_1 and an existing cluster template s_2 is computed as:

$$\text{Sim}(s_1, s_2) = \frac{\sum_{m=1}^{\min(|s_1|, |s_2|)} \mathbb{I}(s_1[m] = s_2[m])}{\max(|s_1|, |s_2|)} \quad (2)$$

where $\mathbb{I}(\cdot)$ is the indicator function. If $\text{Sim}(s_1, s_2) \geq \theta_{\text{sim}}$ (calibrated to $\theta_{\text{sim}} = 0.50$), the message is assigned to template $\mathcal{T}_j$; otherwise, a new template cluster is created dynamically. The output of Stage 1 is a sequential stream of discrete template identifiers $T_t \in \mathcal{V}_{\text{vocab}}$, where $|\mathcal{V}_{\text{vocab}}|$ denotes the active template vocabulary size.

3.3 Hybrid TCN-Transformer Autoencoder

To detect subtle temporal and distributional anomalies in log sequences, we design a hybrid deep autoencoder combining Temporal Convolutional Networks (TCN) [6] with Transformer Self-Attention [7].

Embedding Layer A sliding window of length $L = 20$ (with stride $S = 1$) extracts template sequences $X = [T_1, T_2, \dots, T_L]$. Each template T_i is mapped to a continuous dense representation using a learned embedding matrix $\mathbf{W}_E \in \mathbb{R}^{|\mathcal{V}| \times d_{\text{model}}}$ and augmented with sinusoidal positional encodings $\mathbf{P} \in \mathbb{R}^{L \times d_{\text{model}}}$:

$$\mathbf{E}_0 = \mathbf{X}\mathbf{W}_E + \mathbf{P}, \quad \mathbf{E}_0 \in \mathbb{R}^{L \times d_{\text{model}}} \quad (3)$$

where $d_{\text{model}} = 64$.

TCN Temporal Feature Extractor To extract localized short-term temporal dependencies with low computational overhead, the embedded sequence $\mathbf{E}_0$ passes through 2 dilated causal 1D convolutional layers. A dilated convolution operation $*_d$ on input sequence $\mathbf{x}$ with kernel filter w of size $K = 3$ and dilation factor d is formulated as:

$$y(t) = (\mathbf{x} *_d w)(t) = \sum_{i=0}^{K-1} w(i) \cdot \mathbf{x}(t - d \cdot i) \quad (4)$$

We configure dilation factors $d \in \{1, 2\}$, enabling an exponential receptive field across neighboring log events without causing feature leakage from future tokens. The TCN output is passed through Layer Normalization and GELU activation:

$$\mathbf{H}_{\text{TCN}} = \text{LayerNorm}(\text{GELU}(\text{TCN}(\mathbf{E}_0))) \quad (5)$$

Multi-Head Self-Attention Transformer To capture complex, long-range correlations and permutations across the entire log sequence, $\mathbf{H}_{\text{TCN}}$ is fed into a

2-layer Transformer Encoder. For each attention head $h \in \{1, \dots, H\}$ ($H = 4$), query $\mathbf{Q}$, key $\mathbf{K}$, and value $\mathbf{V}$ projections are computed:

$$\mathbf{Q}_h = \mathbf{H}\mathbf{W}_Q^h, \quad \mathbf{K}_h = \mathbf{H}\mathbf{W}_K^h, \quad \mathbf{V}_h = \mathbf{H}\mathbf{W}_V^h \quad (6)$$

The scaled dot-product attention is calculated as:

$$\text{Attention}(\mathbf{Q}_h, \mathbf{K}_h, \mathbf{V}_h) = \text{softmax}\left(\frac{\mathbf{Q}_h \mathbf{K}_h^T}{\sqrt{d_k}}\right) \mathbf{V}_h \quad (7)$$

where $d_k = d_{\text{model}}/H = 16$. Multi-head outputs are concatenated and projected:

$$\text{MHA}(\mathbf{H}) = \text{Concat}(\text{head}_1, \dots, \text{head}_H) \mathbf{W}_O \quad (8)$$

The encoded representations pass through a Position-Wise Feed-Forward Network (FFN):

$$\mathbf{Z} = \text{LayerNorm}(\mathbf{H} + \text{MHA}(\mathbf{H})) \quad (9)$$

$$\mathbf{H}_{\text{out}} = \text{LayerNorm}(\mathbf{Z} + \text{FFN}(\mathbf{Z})) \quad (10)$$

Sequence Reconstruction and Anomaly Scoring A linear classification head projects $\mathbf{H}_{\text{out}}$ back to the template vocabulary space $\mathbb{R}^{|\mathcal{V}|}$. During training on normal baseline log sequences, the model optimizes the Cross-Entropy reconstruction loss:

$$\mathcal{L}_{\text{recon}} = -\frac{1}{L} \sum_{t=1}^L \log P(\hat{T}_t = T_t \mid X) \quad (11)$$

At inference time, the anomaly score $S(X)$ for a sequence X is computed as the mean reconstruction cross-entropy:

$$S(X) = \frac{1}{L} \sum_{t=1}^L \mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}_t, T_t) \quad (12)$$

A sequence is flagged as anomalous if $S(X) > \tau$, where threshold τ is calibrated from validation set statistics ($\tau = \mu_{\text{val}} + 3\sigma_{\text{val}}$).

3.4 Foundation Model RCA Narration Engine

When an anomalous sequence X_{anom} is detected ($S(X) > \tau$), traditional systems stop at raising a raw alert. In contrast, our framework triggers the Foundation Model Narration Engine to construct an enriched contextual prompt $\mathcal{P}$.

Context-Enriched Prompt Construction The prompt $\mathcal{P}$ integrates three distinct telemetry dimensions:

1. **Sequence Telemetry:** The ordered list of anomalous log templates along with extracted parameters (source IPs, target ports, user roles, return codes).
2. **Infrastructure Metadata:** The originating network segment (e.g., Academic VLAN 10, Server Farm VLAN 50), device hostname, and service criticality level.
3. **Statistical Indicators:** Numerical anomaly score $S(X)$, confidence delta above threshold $\Delta\tau = S(X) - \tau$, and historical event frequency.

Quantized LLM Inference To ensure practical operational deployment on standard enterprise edge gateway servers without requiring enterprise GPU clusters, we deploy quantized open-source foundation models: Microsoft Phi-3 Mini (3.8B parameters) [8] and Mistral-7B [9]. Models are quantized to 4-bit NormalFloat (INT4 / NF4) via AWQ / bitsandbytes [11], reducing active GPU/CPU RAM consumption from 15.2 GB to under 2.4 GB.

The prompt $\mathcal{P}$ is structured to enforce a strict JSON output schema:

```
{
  "severity": "CRITICAL" | "HIGH" | "MEDIUM" | "LOW",
  "root_cause_summary": "<Concise 1-2 sentence description>",
  "technical_explanation": "<Detailed attack vector / fault diagnosis>",
  "remediation_playbook": ["<Step 1>", "<Step 2>", "<Step 3>"]
}
```

4 Experimental Evaluation

4.1 Experimental Setup and Benchmark Datasets

To evaluate the detection fidelity, runtime latency, and explainability of our framework, we conduct extensive experiments on four datasets:

1. **HUFLIT Campus Logs (Real-World Production):** A massive 19.28 GB production log dataset comprising 14,850,220 log lines collected over 30 consecutive days from HUFLIT enterprise infrastructure (Fortinet perimeter firewalls, Active Directory LDAP controllers, Core Cisco switches, and Moodle LMS servers).
2. **BGL Supercomputer** [1]: Blue Gene/L supercomputer log dataset (744 MB, 4,747,963 log lines) containing annotated hardware failures and system alerts.
3. **HDFS Distributed Cluster** [2]: Hadoop Distributed File System log dataset (1.58 GB, 11,175,629 log lines) containing block generation and network timeout anomalies.
4. **UNSW-NB15 Intrusion Benchmark** [14]: Network telemetry dataset (100 MB, 254,004 records) containing diverse attack vectors (DoS, Backdoors, Fuzzers, Reconnaissance).

Table 2 outlines the dataset characteristics and the multi-dataset performance metrics of our framework.

Table 2. Multi-Dataset Benchmark Characteristics and Empirical Performance of the Proposed Framework

Dataset	Domain	Raw Size	Log Lines	Precision	Recall	F1-Score	ROC-AUC	Latency (ms)
HUFLIT Campus Logs	Enterprise Campus Network	19.28 GB	14,850,220	0.9482	0.9561	0.9521	0.9784	0.0128
BGL Supercomputer	Supercomputing Infrastructure	744 MB	4,747,963	0.9320	0.9450	0.9384	0.9620	0.0142
HDFS Distributed Cluster	Distributed Cloud Storage	1.58 GB	11,175,629	0.9410	0.9520	0.9465	0.9710	0.0135
UNSW-NB15 Intrusion	Cyber Security Telemetry	100 MB	254,004	0.9610	0.9680	0.9645	0.9840	0.0098

4.2 Detection Accuracy Comparison with Baselines

We compare our hybrid TCN-Transformer model against four representative log anomaly detection architectures under identical experimental conditions on the 19.28 GB HUFLIT Campus Dataset: (i) **PCA / Invariants** [1]; (ii) **DeepLog (LSTM)** [2]; (iii) **LogAnomaly (FT-LSTM)** [3]; (iv) **Standard Transformer Autoencoder (No TCN)**.

As detailed in Table 3, our proposed TCN-Transformer achieves the highest F1-score (0.9521) and ROC-AUC (0.9784), outperforming DeepLog ($F1 = 0.8983$) and LogAnomaly ($F1 = 0.9123$). The combination of TCN causal convolutions and multi-head self-attention enables effective modeling of both localized token bursts and global sequential permutations.

Table 3. Anomaly Detection Accuracy Comparison on HUFLIT Production Campus Logs

Model Architecture	Precision	Recall	F1-Score	ROC-AUC	Inference Latency
PCA / Invariant Mining [1]	0.7420	0.6850	0.7123	0.7310	0.0495 ms
DeepLog (LSTM) [2]	0.8850	0.9120	0.8983	0.9150	0.0245 ms
LogAnomaly (FT-LSTM) [3]	0.9010	0.9240	0.9123	0.9320	0.0312 ms
Transformer AE (No TCN)	0.9150	0.9280	0.9214	0.9450	0.0185 ms
Proposed TCN-Transformer	0.9482	0.9561	0.9521	0.9784	0.0128 ms

4.3 Loss Convergence and ROC/PR Curve Analysis

Fig. 2 shows the training and validation reconstruction loss convergence of the TCN-Transformer model across 10 training epochs on the HUFLIT dataset. The model converges stably within 6 epochs without overfitting. Fig. 3 presents the Receiver Operating Characteristic (ROC) and Precision-Recall (PR) curves across the HUFLIT, BGL, and HDFS datasets, validating robust generalization across disparate computing domains.

4.4 Ablation Study: Architecture Components & Quantization Impact

To quantify the individual contributions of each design component, we conduct an ablation study analyzing: (1) Model architectural variants; (2) Quantization precision and Foundation Model selection for the RCA engine.




Fig. 2. Reconstruction Loss Convergence.




Fig. 3. ROC and Precision-Recall Curves.

As shown in Table 4, removing the TCN layer degrades F1-score from 0.9521 to 0.9214, demonstrating the necessity of dilated convolutions for localized temporal feature extraction. Furthermore, quantizing the Phi-3 Mini foundation model to INT4 reduces memory consumption by 76.3% (from 7.6 GB to 1.8 GB) while preserving 97.6% of the RCA narration semantic fidelity (ROUGE-L of 0.8420 vs 0.8625 in FP16).

Table 4. Ablation Study on Model Components, Quantization Precision, and RCA Narration Quality

Configuration / Variant	Detection F1	ROC-AUC	ROUGE-L	BLEU-4	LLM Latency	Memory Footprint
TCN Only (No Transformer)	0.8840	0.8990	—	—	—	0.28 MB
Transformer Only (No TCN)	0.9214	0.9450	—	—	—	1.45 MB
Full TCN-Transformer (Detector)	0.9521	0.9784	—	—	—	1.68 MB
Phi-3 Mini (FP16 Baseline)	—	—	0.8625	0.8010	285.40 ms	7.60 GB
Mistral-7B (INT4 Quantized)	—	—	0.8540	0.7950	198.50 ms	4.10 GB
Phi-3 Mini (INT4 Quantized - Ours)	—	—	0.8420	0.7815	125.18 ms	1.80 GB

4.5 Incident Triage Latency & SOC Operational Efficiency

To evaluate operational impact, we measured the end-to-end incident triage time across 50 simulated enterprise campus security incidents. Fig. 4 illustrates the latency comparison between traditional manual SOC triage and our LLM-augmented automated workflow. Manual log analysis, correlation, and remediation drafting required an average of 15.4 minutes (924 seconds) per incident. In contrast, our automated pipeline executes TCN-Transformer detection in 0.0128 ms and completes full LLM RCA narrative synthesis in 125.18 ms, reducing mean incident resolution delay by over 99.98%.

5 Real-World Campus Network Case Studies

To qualitatively illustrate the explainability and operational utility of our framework, we present two representative incident investigations extracted from the HUFLIT production dataset.




Fig. 4. Incident Triage Latency Comparison (Log Scale): Manual Human Expert Inspection (15.4 min) versus LLM-Augmented Automated RCA Narration (125.18 ms).

5.1 Case Study 1: Active Directory Credential Stuffing Attack

Incident Context: A distributed botnet initiated a credential stuffing attack against the campus Active Directory LDAP service. Within 3.2 seconds, 45 authentication failure events were recorded across student Wi-Fi VLAN subnets ($S(X) = 0.942 > \tau$). The Foundation Model generates the following structured diagnosis:

- **Severity:** CRITICAL
- **Root Cause:** High-density Credential Stuffing / Brute-Force Authentication targeting campus LDAP controllers.
- **Technical Diagnosis:** 45 consecutive failed login attempts within 3.2 s from untrusted VLAN subnets indicate automated password dictionary spraying.
- **Remediation Playbook:** (1) Install OpenFlow hard-drop rule on Core Switch for offending source IP/MACs; (2) Enforce 15-minute lockout on targeted accounts; (3) Dispatch out-of-band 2FA re-authentication prompts.

5.2 Case Study 2: Inter-VLAN Spanning Tree Protocol (STP) Loop

Incident Context: An unauthorized switch connected to a dormitory access port created a Layer 2 loop, triggering a broadcast storm that saturated the inter-building optical uplink ($S(X) = 0.887 > \tau$). The Foundation Model produces:

- **Severity:** HIGH
- **Root Cause:** Inter-VLAN Spanning Tree Protocol (STP) loop causing broadcast storm on Edge Switch 4.
- **Technical Diagnosis:** Abnormal surge of broadcast packet log templates with high egress queue drops confirms an unmanaged bridge loop saturating the primary distribution trunk.
- **Remediation Playbook:** (1) Enable BPDU Guard and Root Guard on Port 12 of Edge Switch 4; (2) Re-route administrative VLAN traffic to backup optical link; (3) Issue automated dispatch ticket for field inspection.

6 Discussion & Threats to Validity

Internal Validity: The ground-truth anomaly annotations on the HUFLIT production dataset were curated and independently validated by senior campus network security administrators. The multi-dataset evaluation across BGL, HDFS, and UNSW-NB15 mitigates dataset-specific bias.

External Validity: By evaluating on diverse system topologies (enterprise campus syslog, supercomputer event logs, cloud distributed storage, and packet flow telemetry), we demonstrate the broad applicability of the TCN-Transformer architecture across distinct AIOps domains.

LLM Hallucination Safeguards: To eliminate factual hallucinations in LLM-generated playbooks, the prompt $\mathcal{P}$ uses deterministic temperature scaling ($T = 0.1$) and constrains the output space to verified network actions extracted from institutional standard operating procedures.

7 Conclusion

This paper presented an LLM-Augmented AIOps Framework for real-time log anomaly detection and automated Root Cause Analysis in enterprise campus infrastructures. By synthesizing online Drain3 parsing, a hybrid TCN-Transformer sequence autoencoder, and a lightweight INT4-quantized Foundation Model, the system achieves superior anomaly detection accuracy ($F1 = 0.9521$, ROC-AUC = 0.9784 on 19.28 GB real-world campus logs), sub-millisecond detection inference (0.0128 ms), and reduces human incident triage delay by over 99.9% (125.18 ms automated narration vs 15.4 minutes manual inspection). Future work will explore multi-agent LLM consensus for autonomous closed-loop network self-healing.

Acknowledgment

The authors thank the Faculty of Information Technology, Ho Chi Minh City University of Foreign Languages - Information Technology (HUFLIT), Ho Chi Minh City, Vietnam for providing the production campus log telemetry and research computing infrastructure.

References

1. He, P., Zhuang, J., Ma, X., et al.: Drain: An Online Log Parsing Approach Based on Fixed Depth Tree. In: Proc. IEEE Int. Conf. Web Serv. (ICWS), pp. 33–40 (2017).
2. Du, M., Li, F., Zheng, G., Srikumar, V.: DeepLog: Anomaly Detection and Diagnosis from System Logs Through Deep Learning. In: Proc. ACM SIGSAC Conf. Comput. Commun. Secur. (CCS), pp. 1285–1298 (2017).

3. Meng, W., Liu, Y., Zhu, Y., et al.: LogAnomaly: Unsupervised Detection of Sequential and Quantitative Anomalies in Unstructured Logs. In: Proc. 28th Int. Joint Conf. Artif. Intell. (IJCAI), pp. 4739–4745 (2019).
4. Zhang, X., Xu, Y., Lin, Q., et al.: Robust Log-Based Anomaly Detection on Unstable Log Data. In: Proc. 27th ACM Joint Eur. Softw. Eng. Conf. Symp. Found. Softw. Eng. (ESEC/FSE), pp. 807–817 (2019).
5. Li, X., Chen, P., Jing, L., et al.: SwissLog: Robust and Unified Log Anomaly Detection via Semantic Representation. *IEEE Trans. Serv. Comput.* 15(6), 3326–3338 (2022).
6. Bai, S., Kolter, J.Z., Koltun, V.: An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling. *arXiv preprint arXiv:1803.01271* (2018).
7. Vaswani, A., Shazeer, N., Parmar, N., et al.: Attention Is All You Need. In: *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 30, pp. 5998–6008 (2017).
8. Abdadin, M., et al.: Phi-3 Technical Report: Highly Capable Language Model Locally on Your Phone. *arXiv preprint arXiv:2404.14219* (2024).
9. Jiang, A.Q., Sablayrolles, A., Mensch, A., et al.: Mistral 7B. *arXiv preprint arXiv:2310.06825* (2023).
10. Hu, E.J., Shen, Y., Wallis, P., et al.: LoRA: Low-Rank Adaptation of Large Language Models. In: *Proc. Int. Conf. Learn. Represent. (ICLR)* (2022).
11. Dettmers, T., Pagnoni, A., Holtzman, A., Zettlemoyer, L.: QLoRA: Efficient Fine-tuning of Quantized LLMs. In: *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 36 (2023).
12. Ahmed, T., Ghosh, S., Bansal, C., et al.: Recommending Root-Cause and Mitigation Steps for Cloud Incidents using Large Language Models. In: *Proc. IEEE/ACM Int. Conf. Softw. Eng. (ICSE)*, pp. 1–13 (2023).
13. Soldani, J., Brogi, A.: Anomaly Detection and Failure Root Cause Analysis in (Micro)Service-Based Cloud Applications: A Survey. *ACM Comput. Surv.* 55(3), 1–39 (2022).
14. Moustafa, N., Slay, J.: UNSW-NB15: A Comprehensive Data Set for Network Intrusion Detection Systems. In: *Proc. IEEE Mil. Commun. Inf. Syst. Conf. (MilCIS)*, pp. 1–6 (2015).
15. Du, M., Li, F.: Spell: Streaming Parsing of System Event Logs. In: *Proc. IEEE Int. Conf. Data Min. (ICDM)*, pp. 859–864 (2016).
16. Nedelkoski, S., Bogatinovski, J., Mandapati, A.K., et al.: Multi-Source Distributed System Log Anomaly Detection with Functional Transformers. In: *Proc. ACM SIGKDD Conf. Knowl. Discov. Data Min. (KDD)*, pp. 2487–2495 (2020).
17. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In: *Proc. NAACL-HLT*, pp. 4171–4186 (2019).
18. Chen, B., Zhang, I., et al.: Empirical Analysis of Foundation Models in AIOps: Opportunities and Open Challenges. *IEEE Softw.* 41(2), 45–53 (2024).
19. Rose, S., Borchert, O., Mitchell, S., Connelly, S.: Zero Trust Architecture. *NIST Special Publication 800-207* (2020).
20. Notaro, P., et al.: A Survey on Artificial Intelligence for IT Operations (AIOps). *IEEE Commun. Surv. Tutorials* 26(1), 542–578 (2024).